

image

September 26, 2021

Copyright 2020 Google LLC.

```
[ ]: # Licensed under the Apache License, Version 2.0 (the "License");  
# you may not use this file except in compliance with the License.  
# You may obtain a copy of the License at  
#  
# https://www.apache.org/licenses/LICENSE-2.0  
#  
# Unless required by applicable law or agreed to in writing, software  
# distributed under the License is distributed on an "AS IS" BASIS,  
# WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.  
# See the License for the specific language governing permissions and  
# limitations under the License.
```

1 Introduction to Image Classification

We have learned about binary and multiclass classification, and we've done so using datasets consisting of feature columns that contain numeric and string values. The numbers could be continuous or categorical. The strings we have used so far were all categorical features.

In this lab we will perform another type of classification: **image classification**.

Image classification can be binary: *"Is this an image of a dog?"*

It can also be multiclass: *"Is this an image of a cat, dog, horse, or cow?"*

The questions above assume there is only one item in an image. There is an even more advanced form of multiclass classification that answers the following question: What are all of the classes in an image and where are they located? For example: *"Where are all of the cats, dogs, horses, and cows in this image?"*

In this introduction to image classification, we'll focus on classification where there is only one item depicted in each image. In future labs we'll learn about the more advanced forms of image classification.

1.1 The Dataset

The dataset we'll use for this Colab is the [Fashion-MNIST](#) dataset, which contains 70,000 grayscale images labeled with one of ten categories.

The categories are:

Label	Class
0	T-shirt/top
1	Trouser
2	Pullover
3	Dress
4	Coat
5	Sandal
6	Shirt
7	Sneaker
8	Bag
9	Ankle boot

The images show individual articles of clothing at low resolution (28 by 28 pixels), as seen here:

Figure 1. Fashion-MNIST samples (by Zalando, MIT License).

1.1.1 Load the Data

Now that we have a rough understanding of the data we're going to use in our model, let's load the data into this lab. The Fashion MNIST dataset is conveniently available from the [Keras Datasets repository](#) along with a [utility function](#) for downloading and loading the data into NumPy arrays.

In the code cell below, we import TensorFlow and download the Fashion-MNIST data.

```
[ ]: import tensorflow as tf

(train_images, train_labels), (test_images, test_labels) = \
    tf.keras.datasets.fashion_mnist.load_data()

print(train_images.shape)
print(train_labels.shape)
print(test_images.shape)
print(test_labels.shape)
```

```
(60000, 28, 28)
```

```
(60000,)
```

```
(10000, 28, 28)
```

```
(10000,)
```

`load_data()` returns two tuples, one for the training dataset and the other for the testing dataset. As you can see from the output of the code cell above, we have 60,000 training samples and 10,000 testing samples. This makes for a 14% holdout of the data.

You might be wondering what that 28, 28 is in the image data. That is a two-dimensional representation of the image. This is our feature data. Each pixel of the image is a feature. A 28 by 28 image has 784 pixels.

As you can see, even a tiny image generates quite a few features. If we were processing 4k-resolution images, which are often 3840 by 2160 pixels, then we would have 8,294,400 features! Over eight

million features is quite a bit. In later labs we'll address some strategies for working with this massive amount of data.

1.1.2 Exploratory Data Analysis

It is always a good idea to look at your data before diving in to building your model. Remember that our data is divided across four NumPy arrays, two of which are three-dimensional arrays:

```
[ ]: print('Training images:', train_images.shape)
      print('Training labels:', train_labels.shape)
      print('Test images:', test_images.shape)
      print('Test labels:', test_labels.shape)
```

```
Training images: (60000, 28, 28)
Training labels: (60000,)
Test images: (10000, 28, 28)
Test labels: (10000,)
```

To make our exploration tasks a little easier, let's put the data into a Pandas `DataFrame`. One way to do this is to flatten the 28 by 28 image into a flat array of 784 pixels, with the pixel number being the column name. We then add the labels to a `target` column.

```
[ ]: import numpy as np
      import pandas as pd

      train_df = pd.DataFrame(
          np.array([x.flatten() for x in train_images]),
          columns=[i for i in range(784)]
      )
      train_df['target'] = train_labels

      train_df.describe()
```

```
[ ]:
count    0          1          2          3          4  \
count    60000.000000  60000.000000  60000.000000  60000.000000  60000.000000
mean         0.000800    0.005783    0.030083    0.103800    0.249683
std         0.092554    0.249033    0.767868    2.512017    4.331376
min          0.000000    0.000000    0.000000    0.000000    0.000000
25%          0.000000    0.000000    0.000000    0.000000    0.000000
50%          0.000000    0.000000    0.000000    0.000000    0.000000
75%          0.000000    0.000000    0.000000    0.000000    0.000000
max         16.000000   36.000000   119.000000   164.000000   224.000000

count    5          6          7          8          9  \
count    60000.000000  60000.000000  60000.000000  60000.000000  60000.000000
mean         0.414717    0.821667    2.224733    5.698667   14.434650
std         5.827394    8.309935   14.201820   23.835980   38.204702
min          0.000000    0.000000    0.000000    0.000000    0.000000
25%          0.000000    0.000000    0.000000    0.000000    0.000000
```

50%	0.000000	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000	0.000000	0.000000
max	230.000000	221.000000	221.000000	254.000000	255.000000

	...	775	776	777	778 \
count	...	60000.000000	60000.000000	60000.000000	60000.000000
mean	...	23.208633	16.576250	17.831967	22.918850
std	...	48.881430	42.044318	43.911297	51.928401
min	...	0.000000	0.000000	0.000000	0.000000
25%	...	0.000000	0.000000	0.000000	0.000000
50%	...	0.000000	0.000000	0.000000	0.000000
75%	...	8.000000	0.000000	0.000000	0.000000
max	...	255.000000	255.000000	255.000000	255.000000

		779	780	781	782	783 \
count	60000.000000	60000.000000	60000.000000	60000.000000	60000.000000	60000.000000
mean	17.916900	8.485717	2.706333	0.819000	0.070883	
std	45.173634	29.448614	17.258682	9.133252	2.075829	
min	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	0.000000	0.000000	0.000000	
75%	0.000000	0.000000	0.000000	0.000000	0.000000	
max	255.000000	255.000000	255.000000	255.000000	170.000000	

	target
count	60000.000000
mean	4.500000
std	2.872305
min	0.000000
25%	2.000000
50%	4.500000
75%	7.000000
max	9.000000

[8 rows x 785 columns]

With so many columns, reading the output of `describe()` is nearly impossible. Let's instead do our analysis a little differently.

To begin, we will find the minimum value of every pixel column and output the sorted list of unique values.

```
[ ]: FEATURES = train_df.columns[:-1]

sorted(train_df.loc[:, FEATURES].min().unique())
```

```
[ ]: [0]
```

All of the values were 0.

Let's do the same for the maximum values.

```
[ ]: sorted(train_df.loc[:, FEATURES].max().unique())
```


```
[ ]: [16,  
      36,  
      83,  
      105,  
      119,  
      164,  
      167,  
      170,  
      180,  
      188,  
      189,  
      202,  
      206,  
      211,  
      212,  
      219,  
      221,  
      224,  
      225,  
      227,  
      230,  
      232,  
      233,  
      235,  
      237,  
      239,  
      242,  
      243,  
      244,  
      245,  
      247,  
      249,  
      250,  
      251,  
      252,  
      253,  
      254,  
      255]
```

That is more interesting. We seem to have values ranging from 16 through 255. These values represent color intensities for grayscale images. 0, which we saw as a minimum value, maps to black in the color map that we will use, while 255 is white.

Let's see a histogram distribution of our max pixel values.

```
[ ]: import matplotlib.pyplot as plt

_ = plt.hist(train_df.loc[:, FEATURES].max().unique())
```



Unsurprisingly, higher intensity values seem to be more prevalent as maximum pixel values than lower intensity values.

Exercise 1: Charting Pixel Intensities In the example above, we created a histogram containing the maximum pixel intensities. In this exercise you will create a histogram for all pixel intensities in the training dataset.

If some intensities are outliers, remove them to get a more meaningful histogram.

Hint: The NumPy `where` and `flatten` can come in handy for this exercise.

Student Solution

```
[ ]: train_df.loc[FEATURES]
```

```
[ ]:
      0  1  2  3  4  5  6  7  8  9  ...  775  776  777  778  779  780  \
0      0  0  0  0  0  0  0  0  0  0  ...   0   0   0   0   0   0
1      0  0  0  0  0  1  0  0  0  0  ...  114  130  76   0   0   0
2      0  0  0  0  0  0  0  0  0  0  ...   0   1   0   0   0   0
3      0  0  0  0  0  0  0  0  0  33  ...   0   0   0   0   0   0
4      0  0  0  0  0  0  0  0  0  0  ...   0   0   0   0   0   0
...    ..  ..  ..  ..  ..  ..  ..  ..  ..  ..  ...  ...  ...  ...  ...
779    0  0  0  0  0  0  0  0  0  1  ...   0   0   0  139  163  58
780    0  0  0  0  0  0  0  3  1  0  ...  19   0   0   0   0   0
```


781	0	0	0	0	0	0	0	38	72	75	91	...	0	0	0	0	0	0
782	0	0	0	0	0	0	0	0	0	3	113	...	3	0	0	0	0	0
783	0	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0

	781	782	783	target
0	0	0	0	9
1	0	0	0	0
2	0	0	0	0
3	0	0	0	3
4	0	0	0	0
..
779	0	0	0	6
780	0	0	0	0
781	9	0	0	6
782	0	0	0	1
783	0	0	0	1

[784 rows x 785 columns]

```
[ ]: inten = train_df.loc[:, FEATURES].to_numpy().flatten()
plt.hist(inten)
```

```
[ ]: (array([25790964., 1344502., 1582733., 1643562., 1876736., 2034790.,
2606932., 3303752., 4306392., 2549637.]),
array([ 0. , 25.5, 51. , 76.5, 102. , 127.5, 153. , 178.5, 204. ,
229.5, 255. ]),
<BarContainer object of 10 artists>)
```



Continuing on With EDA Now that we have a basic idea of the values in our dataset, let's see if any are missing.

```
[ ]: train_df.isna().any().any()  
      #train_df.isna().any().any()
```

```
[ ]: False
```


Good. We now know we aren't missing any values, and our pixel values range from 0 through 255. Let's now see if our target values are what we expect.

```
[ ]: sorted(train_df['target'].unique())
```

```
[ ]: [0, 1, 2, 3, 4, 5, 6, 7, 8, 9]
```

Let's see the distribution.

```
[ ]: _ = train_df['target'].hist(bins=[0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10])
```



The class types seem evenly distributed. We have 6,000 of each.

The numeric values should map to these clothing types:

Label	Class
0	T-shirt/top
1	Trouser
2	Pullover
3	Dress
4	Coat
5	Sandal
6	Shirt
7	Sneaker
8	Bag
9	Ankle boot


We can spot check this by looking at some of the images. Let's check a random 'T-shirt/top'.

To do this we select a random index from the 'T-shirt/top' items (`target = 0`). We then reshape the pixel columns back into a 28 by 28 two-dimensional array, which are the dimensions of the image. We then use `imshow()` to display the image.

```
[ ]: index = np.random.choice(train_df[train_df['target'] == 0].index.values)

pixels = train_df.loc[index, FEATURES].to_numpy().reshape(28, 28)

_ = plt.imshow(pixels, cmap='gray')
```



In our sample we got an image that looked like a very low resolution t-shirt. You should see the same. Note: every time you rerun the above cell, a new random index will be chosen, so feel free

to cycle through some of the values to see the different types of t-shirt/top images included in the dataset.

This single image spot checking is okay, but it doesn't scale well.

We can view multiple images at a time using the `GridSpec` class from Matplotlib.

In the code below, we build a visualization with a 10 by 10 grid of images in our t-shirt class.

The code imports `gridspec`, sets the number of rows and columns, and then sets the figure size so the image is large enough for us to actually see different samples.

After that bit of setup, we create a 10 by 10 `GridSpec`. The other parameters to the constructor are there to ensure the images are tightly packed into the grid. Try experimenting with some other values.

Next we randomly choose 100 indexes from items labelled with class 0 our training data.

The remainder of the code should look pretty familiar. We used similar code above to show a single image. The difference in this code is that we are adding 100 subplots using the `GridSpec`.

Label	Class
0	T-shirt/top
1	Trouser
2	Pullover
3	Dress
4	Coat
5	Sandal
6	Shirt
7	Sneaker
8	Bag
9	Ankle boot

```
[ ]:
[ ]: from matplotlib import gridspec

# Row and column count (100 samples)
rows = 10
cols = 10

# Size of the final output image
plt.figure(figsize=(12, 12))

# Grid that will be used to organize our samples
gspec = gridspec.GridSpec(
    rows,
    cols,
    wspace = 0.0,
    hspace = 0.0,
```

```

    top = 1.0,
    bottom = 0.0,
    left = 0.00,
    right = 1.0,
)


# Randomly choose a sample of t-shirts
T_SHIRTS = 0
indexes = np.random.choice(
    train_df[train_df['target'] == T_SHIRTS].index.values,
    rows*cols
)

# Add each sample to a plot using the GridSpec
cnt = 0
for r in range(rows):
    for c in range(cols):
        row = train_df.loc[indexes[cnt], FEATURES]
        img = row.to_numpy().reshape((28, 28))

        ax = plt.subplot(gspec[r, c])
        ax.imshow(img, cmap='gray')
        ax.xaxis.set_visible(False)
        ax.yaxis.set_visible(False)
        cnt = cnt + 1

plt.show()

```



Exercise 2: Visualizing Every Class In this exercise, you'll take the code that we used above to visualize t-shirts and use it to visualize every class represented in our dataset. You'll need to print out the class name and then show a 10 by 10 grid of samples from that class. Try to minimize the amount of repeated code in your solution.

Student Solution Defining a function to pass in a particular class

```
[ ]: def function(classer):
    rows = 10
    cols = 10

    # Size of the final output image
```

```

plt.figure(figsize=(12, 12))

# Grid that will be used to organize our samples
gspec = gridspec.GridSpec(
    rows,
    cols,
    wspace = 0.0,
    hspace = 0.0,
    top = 1.0,
    bottom = 0.0,
    left = 0.00,
    right = 1.0,
)

# Randomly choose a sample of t-shirts

indexes = np.random.choice(
    train_df[train_df['target'] == classer].index.values,
    rows*cols
)
cnt = 0
for r in range(rows):
    for c in range(cols):
        row = train_df.loc[indexes[cnt], FEATURES]
        img = row.to_numpy().reshape((28, 28))
        ax = plt.subplot(gspec[r, c])
        ax.imshow(img, cmap='gray')
        ax.xaxis.set_visible(False)
        ax.yaxis.set_visible(False)
        cnt = cnt + 1


plt.show()


```


```


[ ]: classes = [0,1,2,3,4,5,6,7]
for i in classes:
    function(i)


```


















Wrapping Up EDA From our visual analysis, our samples seem reasonable.

First off, class names seem to match pictures. This can give us some confidence that our data is labelled correctly.

Another nice thing is that all of the clothing items seem to be oriented in the same direction for the most part. If shoes were pointing in different directions, or if any images were rotated, then we would have had a lot more processing to do.

And finally, all of our images are the same dimensions and are encoded with a single numeric grayscale intensity. In the real world, you'll likely not get so lucky. Images are acquired in different sizes and with different color encodings. We'll get to some examples of this in future labs.

Based on our analysis so far, we can end our EDA and move on to model building.

1.2 Modeling

We have many options for building a multiclass classification model for images. In this lab we will build a deep neural network using TensorFlow Keras.

1.2.1 Preparing the Data

Our feature data is on a scale from 0 to 255, and our target data is categorically encoded. Fortunately, all of the features are on the same scale, so we don't have to worry about standardizing scale. However, we'll need to do a little data preprocessing in order to get our data ready for modeling.

The first bit of data preprocessing we'll do is bring the feature values into the range of 0.0 and 1.0. We could perform normalization to do this, but normalization actually isn't the only solution in this case.

We know that all of our features are pixel values in the range of 0 to 255. We also know from our EDA that every feature has a minimum value of 0, but that the max values have a pretty wide range. It is possible we would make our model worse by normalizing, since we'd be making the same values across pixels not map to the same color.

Instead of normalizing, we can just divide every feature by 255.0. This keeps the relative values the same across pixels.

```
[ ]: train_df[FEATURES] = train_df[FEATURES] / 255.0

train_df[FEATURES].describe()
```

```
[ ]:
```

	0	1	2	3	4	\
count	60000.000000	60000.000000	60000.000000	60000.000000	60000.000000	
mean	0.000003	0.000023	0.000118	0.000407	0.000979	
std	0.000363	0.000977	0.003011	0.009851	0.016986	
min	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	0.000000	0.000000	0.000000	
75%	0.000000	0.000000	0.000000	0.000000	0.000000	
max	0.062745	0.141176	0.466667	0.643137	0.878431	

	5	6	7	8	9	\
count	60000.000000	60000.000000	60000.000000	60000.000000	60000.000000	
mean	0.001626	0.003222	0.008724	0.022348	0.056606	
std	0.022853	0.032588	0.055693	0.093474	0.149822	
min	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	0.000000	0.000000	0.000000	
75%	0.000000	0.000000	0.000000	0.000000	0.000000	
max	0.901961	0.866667	0.866667	0.996078	1.000000	

	...	774	775	776	777	\
count	...	60000.000000	60000.000000	60000.000000	60000.000000	
mean	...	0.135547	0.091014	0.065005	0.069929	
std	...	0.225717	0.191692	0.164880	0.172201	
min	...	0.000000	0.000000	0.000000	0.000000	
25%	...	0.000000	0.000000	0.000000	0.000000	
50%	...	0.000000	0.000000	0.000000	0.000000	
75%	...	0.223529	0.031373	0.000000	0.000000	
max	...	1.000000	1.000000	1.000000	1.000000	

	778	779	780	781	782	\
count	60000.000000	60000.000000	60000.000000	60000.000000	60000.000000	
mean	0.089878	0.070262	0.033277	0.010613	0.003212	
std	0.203641	0.177152	0.115485	0.067681	0.035817	
min	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	0.000000	0.000000	0.000000	
75%	0.000000	0.000000	0.000000	0.000000	0.000000	
max	1.000000	1.000000	1.000000	1.000000	1.000000	

	783
count	60000.000000
mean	0.000278
std	0.008141
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	0.666667

[8 rows x 784 columns]

Exercise 3: One-Hot Encoding Our target values are categorical values in a column named `target`. In this exercise, you will one-hot encode the target values. Your code should:

1. Create ten new columns named `target_0` through `target_9`.
2. Create a variable called `TARGETS` that contains the 10 target column names.
3. `describe()` the ten new target column values to ensure that they have values between 0 and 1 and that the one-hot encoding looks evenly distributed.

Student Solution Getting one hot encoding dataframe

```
[ ]: from sklearn.preprocessing import OneHotEncoder
```

```

target_columns =
    ↳ ['target_0', 'target_1', 'target_2', 'target_3', 'target_4', 'target_5', 'target_6', 'target_7', 't
TARGETS = target_columns
key = ['target']
ohe = OneHotEncoder(sparse=False)

encodedTargets = ohe.fit_transform(train_df[['target']])
print(type(encodedTargets))

encodedTargets_df = pd.DataFrame(encodedTargets)
encodedTargets_df.columns =target_columns

```

```
<class 'numpy.ndarray'>
```

Dropping Target Column

```
[ ]: train_df = train_df.drop('target',axis =1 )
```

```
train_df
```

```
[ ]:
```

	0	1	2	3	4	5	6	7	8	9	...	\
0	0.0	0.0	0.0	0.0	0.0	0.000000	0.0	0.0	0.000000	0.000000	...	
1	0.0	0.0	0.0	0.0	0.0	0.003922	0.0	0.0	0.000000	0.000000	...	
2	0.0	0.0	0.0	0.0	0.0	0.000000	0.0	0.0	0.000000	0.086275	...	
3	0.0	0.0	0.0	0.0	0.0	0.000000	0.0	0.0	0.129412	0.376471	...	
4	0.0	0.0	0.0	0.0	0.0	0.000000	0.0	0.0	0.000000	0.000000	...	
...	
59995	0.0	0.0	0.0	0.0	0.0	0.000000	0.0	0.0	0.000000	0.000000	...	
59996	0.0	0.0	0.0	0.0	0.0	0.000000	0.0	0.0	0.000000	0.000000	...	
59997	0.0	0.0	0.0	0.0	0.0	0.000000	0.0	0.0	0.000000	0.019608	...	
59998	0.0	0.0	0.0	0.0	0.0	0.000000	0.0	0.0	0.000000	0.000000	...	
59999	0.0	0.0	0.0	0.0	0.0	0.000000	0.0	0.0	0.000000	0.000000	...	
	774	775	776	777	778	779	780	781	782	...	\	
0	0.000000	0.000000	0.000000	0.000000	0.0	0.000000	0.0	0.0	0.0			
1	0.466667	0.447059	0.509804	0.298039	0.0	0.000000	0.0	0.0	0.0			
2	0.000000	0.000000	0.003922	0.000000	0.0	0.000000	0.0	0.0	0.0			
3	0.000000	0.000000	0.000000	0.000000	0.0	0.000000	0.0	0.0	0.0			
4	0.000000	0.000000	0.000000	0.000000	0.0	0.000000	0.0	0.0	0.0			
...			
59995	0.000000	0.000000	0.000000	0.000000	0.0	0.000000	0.0	0.0	0.0			
59996	0.000000	0.000000	0.000000	0.000000	0.0	0.000000	0.0	0.0	0.0			
59997	0.000000	0.000000	0.000000	0.000000	0.0	0.000000	0.0	0.0	0.0			
59998	0.258824	0.211765	0.196078	0.019608	0.0	0.003922	0.0	0.0	0.0			
59999	0.000000	0.000000	0.000000	0.000000	0.0	0.000000	0.0	0.0	0.0			


```

783
0      0.0
1      0.0
2      0.0
3      0.0
4      0.0
...
59995  0.0
59996  0.0
59997  0.0
59998  0.0
59999  0.0

```

[60000 rows x 784 columns]

Combining DataFrames

```

[ ]: encodedTargets_df
train_df
frames = [train_df,encodedTargets_df]

train_df = train_df.join(encodedTargets_df)

train_df
#new_target_df = pd.merge(train_df,encodedTargets_df)

```

```

[ ]:
   0      1      2      3      4      5      6      7      8      9  ...  \
0   0.0  0.0  0.0  0.0  0.0  0.000000  0.0  0.0  0.000000  0.000000  ...
1   0.0  0.0  0.0  0.0  0.0  0.003922  0.0  0.0  0.000000  0.000000  ...
2   0.0  0.0  0.0  0.0  0.0  0.000000  0.0  0.0  0.000000  0.086275  ...
3   0.0  0.0  0.0  0.0  0.0  0.000000  0.0  0.0  0.129412  0.376471  ...
4   0.0  0.0  0.0  0.0  0.0  0.000000  0.0  0.0  0.000000  0.000000  ...
...
59995  0.0  0.0  0.0  0.0  0.0  0.000000  0.0  0.0  0.000000  0.000000  ...
59996  0.0  0.0  0.0  0.0  0.0  0.000000  0.0  0.0  0.000000  0.000000  ...
59997  0.0  0.0  0.0  0.0  0.0  0.000000  0.0  0.0  0.000000  0.019608  ...
59998  0.0  0.0  0.0  0.0  0.0  0.000000  0.0  0.0  0.000000  0.000000  ...
59999  0.0  0.0  0.0  0.0  0.0  0.000000  0.0  0.0  0.000000  0.000000  ...

   target_0  target_1  target_2  target_3  target_4  target_5  target_6  \
0         0.0         0.0         0.0         0.0         0.0         0.0
1         1.0         0.0         0.0         0.0         0.0         0.0
2         1.0         0.0         0.0         0.0         0.0         0.0
3         0.0         0.0         0.0         1.0         0.0         0.0
4         1.0         0.0         0.0         0.0         0.0         0.0
...         ...         ...         ...         ...         ...         ...

```

59995	0.0	0.0	0.0	0.0	0.0	1.0	0.0
59996	0.0	1.0	0.0	0.0	0.0	0.0	0.0
59997	0.0	0.0	0.0	1.0	0.0	0.0	0.0
59998	1.0	0.0	0.0	0.0	0.0	0.0	0.0
59999	0.0	0.0	0.0	0.0	0.0	1.0	0.0

	target_7	target_8	target_9
0	0.0	0.0	1.0
1	0.0	0.0	0.0
2	0.0	0.0	0.0
3	0.0	0.0	0.0
4	0.0	0.0	0.0
...
59995	0.0	0.0	0.0
59996	0.0	0.0	0.0
59997	0.0	0.0	0.0
59998	0.0	0.0	0.0
59999	0.0	0.0	0.0

[60000 rows x 794 columns]

1.2.2 Configure and Compile the Model

We'll be relying on the TensorFlow Keras [Sequential model](#) and [Dense layers](#) that we used in previous labs.

In this case our input shape needs to be the size of our feature count. We'll then add a few hidden layers and then use a softmax layer the same width as our target count. This layer should output the probability that a given set of input features maps to each of our targets. The sum of the probabilities will equal 1.0.

```
[ ]: model = tf.keras.Sequential([
    tf.keras.layers.Dense(128, input_shape=(len(FEATURES),)),
    tf.keras.layers.Dense(64, activation=tf.nn.relu),
    tf.keras.layers.Dense(32, activation=tf.nn.relu),
    tf.keras.layers.Dense(len(TARGETS), activation=tf.nn.softmax)
])

model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 128)	100480
dense_1 (Dense)	(None, 64)	8256
dense_2 (Dense)	(None, 32)	2080

```

-----
dense_3 (Dense)                (None, 10)                330
=====

```

```

Total params: 111,146
Trainable params: 111,146
Non-trainable params: 0

```

```

-----
2021-09-26 13:10:07.675669: I tensorflow/core/platform/cpu_feature_guard.cc:142]
This TensorFlow binary is optimized with oneAPI Deep Neural Network Library
(oneDNN) to use the following CPU instructions in performance-critical
operations:  SSE4.1 SSE4.2
To enable them in other operations, rebuild TensorFlow with the appropriate
compiler flags.

```

Note that our images are actually 28 by 28 images. We flattened the images when we loaded them into a dataframe for EDA. However, flattening outside of the model isn't necessary. TensorFlow works in many dimensions. If we wanted to keep our images as 28 by 28 matrices, we could have added a [Flatten](#) layer as shown below.

```

model = tf.keras.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)),
    tf.keras.layers.Dense(128, activation=tf.nn.relu),
    tf.keras.layers.Dense(64, activation=tf.nn.relu),
    tf.keras.layers.Dense(32, activation=tf.nn.relu),
    tf.keras.layers.Dense(len(TARGETS), activation=tf.nn.softmax)
])

```

```

model.summary()

```

```

Model: "sequential_8"

```

```

-----
Layer (type)                Output Shape                Param #
=====
flatten_5 (Flatten)         (None, 784)                 0
-----
dense_28 (Dense)             (None, 128)                 100480
-----
dense_29 (Dense)             (None, 64)                  8256
-----
dense_30 (Dense)             (None, 32)                  2080
-----
dense_31 (Dense)             (None, 10)                  330
=====

```

```

Total params: 111,146
Trainable params: 111,146
Non-trainable params: 0

```

This model results in the same number of trainable parameters as our pre-flattened model; it just

saves us from having to flatten the images outside of TensorFlow.

Before the model is ready for training, it needs a few more settings. These are added during the model's *compile* step:

- *Loss function* — This measures how well the model is doing during training. We want to minimize this function to “steer” the model in the right direction. A large loss would indicate the model is performing poorly in classification tasks, meaning it is not matching input images to the correct class names. (It might classify a boot as a coat, for example.)
- *Optimizer* — This is how the model is updated based on the data it sees and its loss function.
- *Metrics* — This is used to monitor the training and testing steps. The following example uses *accuracy*, the fraction of the images that are correctly classified.

```
[ ]: model.compile(
    loss='categorical_crossentropy',
    optimizer='Adam',
    metrics=['accuracy'],
)

model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 128)	100480
dense_1 (Dense)	(None, 64)	8256
dense_2 (Dense)	(None, 32)	2080
dense_3 (Dense)	(None, 10)	330
Total params: 111,146		
Trainable params: 111,146		
Non-trainable params: 0		

1.3 Train the Model

Training a Keras API neural network model to classify images looks just like all of the other Keras models we have worked with so far. We call this the `model.fit` method, passing it our training data and any other parameters we'd like to use.

```
[ ]: callback = tf.keras.callbacks.EarlyStopping(monitor='loss', patience=5)

history = model.fit(
    train_df[FEATURES],
    train_df[TARGETS],
```

```
epochs=500,  
callbacks=[callback]  
)
```

```
Epoch 1/500  
1875/1875 [=====] - 2s 1ms/step - loss: 0.1560 -  
accuracy: 0.9416  
Epoch 2/500  
1875/1875 [=====] - 2s 1ms/step - loss: 0.1503 -  
accuracy: 0.9427  
Epoch 3/500  
1875/1875 [=====] - 2s 1ms/step - loss: 0.1543 -  
accuracy: 0.9427  
Epoch 4/500  
1875/1875 [=====] - 2s 1ms/step - loss: 0.1499 -  
accuracy: 0.9423  
Epoch 5/500  
1875/1875 [=====] - 2s 993us/step - loss: 0.1462 -  
accuracy: 0.9446  
Epoch 6/500  
1875/1875 [=====] - 2s 1ms/step - loss: 0.1504 -  
accuracy: 0.9435  
Epoch 7/500  
1875/1875 [=====] - 2s 1ms/step - loss: 0.1481 -  
accuracy: 0.9440  
Epoch 8/500  
1875/1875 [=====] - 2s 1ms/step - loss: 0.1468 -  
accuracy: 0.9445  
Epoch 9/500  
1875/1875 [=====] - 2s 1ms/step - loss: 0.1473 -  
accuracy: 0.9441  
Epoch 10/500  
1875/1875 [=====] - 2s 1ms/step - loss: 0.1445 -  
accuracy: 0.9445  
Epoch 11/500  
1875/1875 [=====] - 2s 1ms/step - loss: 0.1436 -  
accuracy: 0.9453  
Epoch 12/500  
801/1875 [=====>...] - ETA: 1s - loss: 0.1356 - accuracy:  
0.9488
```

```
-----  
KeyboardInterrupt                                Traceback (most recent call last)  
/var/folders/q7/wrxzkb515gqcskvhd38dw6h0000gn/T/ipykernel_72421/2811613906.py  
↳ in <module>  
    1 callback = tf.keras.callbacks.EarlyStopping(monitor='loss', patience=5)  
    2
```

```

----> 3 history = model.fit(
      4     train_df[FEATURES],
      5     train_df[TARGETS],

~/opt/anaconda3/envs/data/lib/python3.9/site-packages/tensorflow/python/keras/
↳engine/training.py in fit(self, x, y, batch_size, epochs, verbose, callbacks,
↳validation_split, validation_data, shuffle, class_weight, sample_weight,
↳initial_epoch, steps_per_epoch, validation_steps, validation_batch_size,
↳validation_freq, max_queue_size, workers, use_multiprocessing)
    1098         _r=1):
    1099             callbacks.on_train_batch_begin(step)
-> 1100             tmp_logs = self.train_function(iterator)
    1101             if data_handler.should_sync:
    1102                 context.async_wait()

~/opt/anaconda3/envs/data/lib/python3.9/site-packages/tensorflow/python/eager/
↳def_function.py in __call__(self, *args, **kwds)
    826         tracing_count = self.experimental_get_tracing_count()
    827         with trace.Trace(self._name) as tm:
-> 828             result = self._call(*args, **kwds)
    829             compiler = "xla" if self._experimental_compile else "nonXla"
    830             new_tracing_count = self.experimental_get_tracing_count()

~/opt/anaconda3/envs/data/lib/python3.9/site-packages/tensorflow/python/eager/
↳def_function.py in _call(self, *args, **kwds)
    853         # In this case we have created variables on the first call, so we
↳run the
    854         # defunned version which is guaranteed to never create variables.
-> 855         return self._stateless_fn(*args, **kwds) # pylint:
↳disable=not-callable
    856         elif self._stateful_fn is not None:
    857             # Release the lock early so that multiple threads can perform the
↳call

~/opt/anaconda3/envs/data/lib/python3.9/site-packages/tensorflow/python/eager/
↳function.py in __call__(self, *args, **kwargs)
    2940         (graph_function,
    2941          filtered_flat_args) = self._maybe_define_function(args, kwargs)
-> 2942         return graph_function._call_flat(
    2943             filtered_flat_args, captured_inputs=graph_function.
↳captured_inputs) # pylint: disable=protected-access
    2944

~/opt/anaconda3/envs/data/lib/python3.9/site-packages/tensorflow/python/eager/
↳function.py in _call_flat(self, args, captured_inputs, cancellation_manager)
    1916         and executing_eagerly):
    1917             # No tape is watching; skip to running the function.

```

```

-> 1918         return self._build_call_outputs(self._inference_function.call(
1919             ctx, args, cancellation_manager=cancellation_manager))
1920         forward_backward = self._select_forward_and_backward_functions(

~/opt/anaconda3/envs/data/lib/python3.9/site-packages/tensorflow/python/eager/
->function.py in call(self, ctx, args, cancellation_manager)
553         with _InterpolateFunctionError(self):
554             if cancellation_manager is None:
--> 555                 outputs = execute.execute(

556                     str(self.signature.name),
557                     num_outputs=self._num_outputs,

~/opt/anaconda3/envs/data/lib/python3.9/site-packages/tensorflow/python/eager/
->execute.py in quick_execute(op_name, num_outputs, inputs, attrs, ctx, name)
57     try:
58         ctx.ensure_initialized()
---> 59         tensors = pywrap_tfe.TFE_Py_Execute(ctx._handle, device_name, op_
->op_name,

60                                     inputs, attrs, num_outputs)
61     except core._NotOkStatusException as e:

KeyboardInterrupt:

```

As the model trains, the loss and accuracy metrics are displayed. We also store the progression in history.

You'll notice that this took longer to train per epoch than many of the models we've built previously in this course. That's because of the large number of features. Even with these tiny 28 by 28 grayscale images, we are still dealing with 784 features. This is orders of magnitude larger than the 10 or so features we are used to using.

1.3.1 Exercise 4: Graph Model Progress

In this exercise you'll create two graphs. The first will show the model loss over each epoch. The second will show the model accuracy over each epoch. Feel free to use any graphical toolkit we have used so far.

```

[ ]: import matplotlib.pyplot as plt


plt.figure(figsize=(16,5))

plt.subplot(1,2,2)
plt.plot(history.history['loss'])
plt.title('Training Loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train_loss'], loc='best')

```

```
plt.subplot(1,2,1)
plt.plot(history.history['accuracy'])
plt.title('Training Accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train_accuracy'], loc='best')
```

```
[ ]: <matplotlib.legend.Legend at 0x7f95ace62dd0>
```



1.4 Evaluate the Model

Now that our model is trained, let's evaluate it using an independent test data set. Then let's see if the model quality holds up. We'll use `model.evaluate()` and pass in the test dataset. `model.evaluate()` returns a `test_loss` and `test_accuracy`.

Also note that we need to apply the same feature preprocessing to the test data that we did to the train data.

```
[ ]: test_df = pd.DataFrame(
    np.array([x.flatten() for x in test_images]),
    columns=[i for i in range(784)]
)
test_df['target'] = test_labels

test_df[FEATURES] = test_df[FEATURES] / 255.0

for class_i in sorted(test_df['target'].unique()):
    column_name = f'target_{class_i}'
    test_df[column_name] = (test_df['target'] == class_i).astype(int)

train_loss = history.history['loss'][-1]
```



```

train_accuracy = history.history['accuracy'][-1]
(test_loss, test_accuracy) = model.evaluate(test_df[FEATURES], test_df[TARGETS])

print('Training loss:', train_loss)
print('Training accuracy:', train_accuracy)

print('Test loss:', test_loss)
print('Test accuracy:', test_accuracy)

```

```

WARNING:tensorflow:Falling back from v2 loop because of error: Failed to find
data adapter that can handle input: <class 'pandas.core.frame.DataFrame'>,
<class 'NoneType'>
10000/10000 [=====] - 1s 65us/sample - loss: 0.5598 -
accuracy: 0.8761
Training loss: 0.12142243753870328
Training accuracy: 0.9551833
Test loss: 0.5598000788092613
Test accuracy: 0.8761

```

The accuracy on the test dataset is noticeably less than the accuracy on the training dataset. This gap between training accuracy and test accuracy is an example of **overfitting**. Overfitting is when a machine learning model tends to perform worse on new data than on the training data. The trained model is unable to **generalize** to data that it has not seen before.

There are many ways to try to reduce overfitting. One that we have seen is **early stopping**. This causes training to stop when loss stops changing significantly for a model. Without early stopping, the model would continue training, becoming more and more tailored to the training data and likely less to be able to generalize across new data.

Another method for reducing overfitting in deep neural networks is **dropout**. A dropout layer is a layer that sits between two regular layers (in our case, dense layers) and randomly sets some of the values passed between layers to 0.

In TensorFlow the [Dropout](#) class is capable of doing this. To use Dropout you simply add a Dropout layer between other layers of the model. Each dropout layer has a percentage of values that it will set to 0.

```

model = tf.keras.Sequential([
    ...
    tf.keras.layers.Dense(name='L14', 431, activation=tf.nn.relu),
    # Randomly sets 15% of values between L14 and L15 to 0
    tf.keras.layers.Dropout(rate=0.15),
    tf.keras.layers.Dense(name='L15', 257, activation=tf.nn.relu),
    # Randomly sets 1% of values between L15 and L16 to 0
    tf.keras.layers.Dropout(rate=0.01),
    tf.keras.layers.Dense(name='L16', 57, activation=tf.nn.relu),
    ...
])

```

1.4.1 Exercise 5: Dropout Layers

In this exercise take the model from above and add a [Dropout](#) layer or layers between the `Dense` layers. See if you can find a configuration that reduces the gap between the training loss and accuracy and the test loss and accuracy. Document your findings.

Student Solution

```
[ ]: model = tf.keras.Sequential([
    tf.keras.layers.Dense(128, input_shape=(len(FEATURES),)),
    tf.keras.layers.Dropout(rate=0.25),
    tf.keras.layers.Dense(64, activation=tf.nn.relu),
    tf.keras.layers.Dropout(rate=0.15),
    tf.keras.layers.Dense(32, activation=tf.nn.relu),
    tf.keras.layers.Dropout(rate=0.01),
    tf.keras.layers.Dense(len(TARGETS), activation=tf.nn.softmax)
])

model.summary()

model.compile(
    loss='categorical_crossentropy',
    optimizer='Adam',
    metrics=['accuracy'],
)

model.summary()

callback = tf.keras.callbacks.EarlyStopping(monitor='loss', patience=5)

history = model.fit(
    train_df[FEATURES],
    train_df[TARGETS],
    epochs=500,
    callbacks=[callback]
)
```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
dense_16 (Dense)	(None, 128)	100480
dropout (Dropout)	(None, 128)	0
dense_17 (Dense)	(None, 64)	8256
dropout_1 (Dropout)	(None, 64)	0

dense_18 (Dense)	(None, 32)	2080
------------------	------------	------

dropout_2 (Dropout)	(None, 32)	0
---------------------	------------	---

dense_19 (Dense)	(None, 10)	330
------------------	------------	-----

=====
Total params: 111,146
Trainable params: 111,146
Non-trainable params: 0

Model: "sequential_4"

Layer (type)	Output Shape	Param #
--------------	--------------	---------

dense_16 (Dense)	(None, 128)	100480
------------------	-------------	--------

dropout (Dropout)	(None, 128)	0
-------------------	-------------	---

dense_17 (Dense)	(None, 64)	8256
------------------	------------	------

dropout_1 (Dropout)	(None, 64)	0
---------------------	------------	---

dense_18 (Dense)	(None, 32)	2080
------------------	------------	------

dropout_2 (Dropout)	(None, 32)	0
---------------------	------------	---

dense_19 (Dense)	(None, 10)	330
------------------	------------	-----

=====
Total params: 111,146
Trainable params: 111,146
Non-trainable params: 0

WARNING:tensorflow:Falling back from v2 loop because of error: Failed to find data adapter that can handle input: <class 'pandas.core.frame.DataFrame'>, <class 'NoneType'>

Train on 60000 samples

Epoch 1/500

60000/60000 [=====] - 12s 196us/sample - loss: 0.6168 - accuracy: 0.7790

Epoch 2/500

60000/60000 [=====] - 11s 186us/sample - loss: 0.4657 - accuracy: 0.8296

Epoch 3/500

60000/60000 [=====] - 12s 193us/sample - loss: 0.4312 - accuracy: 0.8445

Epoch 4/500

60000/60000 [=====] - 12s 195us/sample - loss: 0.4138 - accuracy: 0.8480

Epoch 5/500
60000/60000 [=====] - 12s 193us/sample - loss: 0.3985 -
accuracy: 0.8544

Epoch 6/500
60000/60000 [=====] - 11s 188us/sample - loss: 0.3900 -
accuracy: 0.8561

Epoch 7/500
60000/60000 [=====] - 11s 186us/sample - loss: 0.3781 -
accuracy: 0.8606

Epoch 8/500
60000/60000 [=====] - 11s 187us/sample - loss: 0.3724 -
accuracy: 0.8619

Epoch 9/500
60000/60000 [=====] - 11s 188us/sample - loss: 0.3690 -
accuracy: 0.8649

Epoch 10/500
60000/60000 [=====] - 11s 187us/sample - loss: 0.3579 -
accuracy: 0.8695

Epoch 11/500
60000/60000 [=====] - 11s 189us/sample - loss: 0.3563 -
accuracy: 0.8683

Epoch 12/500
60000/60000 [=====] - 11s 188us/sample - loss: 0.3542 -
accuracy: 0.8688

Epoch 13/500
60000/60000 [=====] - 11s 189us/sample - loss: 0.3520 -
accuracy: 0.8713

Epoch 14/500
60000/60000 [=====] - 11s 190us/sample - loss: 0.3474 -
accuracy: 0.8732

Epoch 15/500
60000/60000 [=====] - 12s 197us/sample - loss: 0.3420 -
accuracy: 0.8743

Epoch 16/500
60000/60000 [=====] - 12s 199us/sample - loss: 0.3398 -
accuracy: 0.8739

Epoch 17/500
60000/60000 [=====] - 12s 201us/sample - loss: 0.3375 -
accuracy: 0.8744

Epoch 18/500
60000/60000 [=====] - 11s 191us/sample - loss: 0.3355 -
accuracy: 0.8762

Epoch 19/500
60000/60000 [=====] - 11s 190us/sample - loss: 0.3334 -
accuracy: 0.8757

Epoch 20/500
60000/60000 [=====] - 12s 197us/sample - loss: 0.3314 -
accuracy: 0.8779

Epoch 21/500
60000/60000 [=====] - 11s 189us/sample - loss: 0.3280 -
accuracy: 0.8776

Epoch 22/500
60000/60000 [=====] - 12s 198us/sample - loss: 0.3262 -
accuracy: 0.8784

Epoch 23/500
60000/60000 [=====] - 12s 194us/sample - loss: 0.3272 -
accuracy: 0.8781

Epoch 24/500
60000/60000 [=====] - 11s 191us/sample - loss: 0.3222 -
accuracy: 0.8809

Epoch 25/500
60000/60000 [=====] - 12s 200us/sample - loss: 0.3218 -
accuracy: 0.8807

Epoch 26/500
60000/60000 [=====] - 12s 198us/sample - loss: 0.3198 -
accuracy: 0.8815

Epoch 27/500
60000/60000 [=====] - 11s 187us/sample - loss: 0.3186 -
accuracy: 0.8817

Epoch 28/500
60000/60000 [=====] - 11s 188us/sample - loss: 0.3188 -
accuracy: 0.8815

Epoch 29/500
60000/60000 [=====] - 11s 190us/sample - loss: 0.3134 -
accuracy: 0.8842

Epoch 30/500
60000/60000 [=====] - 11s 190us/sample - loss: 0.3148 -
accuracy: 0.8844

Epoch 31/500
60000/60000 [=====] - 11s 187us/sample - loss: 0.3146 -
accuracy: 0.8833

Epoch 32/500
60000/60000 [=====] - 12s 192us/sample - loss: 0.3119 -
accuracy: 0.8839

Epoch 33/500
60000/60000 [=====] - 12s 208us/sample - loss: 0.3111 -
accuracy: 0.8853

Epoch 34/500
60000/60000 [=====] - 12s 199us/sample - loss: 0.3120 -
accuracy: 0.8852

Epoch 35/500
60000/60000 [=====] - 12s 204us/sample - loss: 0.3116 -
accuracy: 0.8834

Epoch 36/500
60000/60000 [=====] - 11s 185us/sample - loss: 0.3057 -
accuracy: 0.8862

Epoch 37/500
 60000/60000 [=====] - 11s 185us/sample - loss: 0.3107 -
 accuracy: 0.8863
 Epoch 38/500
 60000/60000 [=====] - 12s 204us/sample - loss: 0.3059 -
 accuracy: 0.8859
 Epoch 39/500
 60000/60000 [=====] - 12s 193us/sample - loss: 0.3075 -
 accuracy: 0.8859
 Epoch 40/500
 60000/60000 [=====] - 12s 194us/sample - loss: 0.3065 -
 accuracy: 0.8868
 Epoch 41/500
 60000/60000 [=====] - 12s 200us/sample - loss: 0.3046 -
 accuracy: 0.8867
 Epoch 42/500
 60000/60000 [=====] - 12s 205us/sample - loss: 0.3045 -
 accuracy: 0.8870
 Epoch 43/500
 60000/60000 [=====] - 11s 184us/sample - loss: 0.2986 -
 accuracy: 0.8887
 Epoch 44/500
 60000/60000 [=====] - 13s 211us/sample - loss: 0.2999 -
 accuracy: 0.8895
 Epoch 45/500
 60000/60000 [=====] - 11s 190us/sample - loss: 0.3026 -
 accuracy: 0.8872
 Epoch 46/500
 60000/60000 [=====] - 11s 186us/sample - loss: 0.3021 -
 accuracy: 0.8878
 Epoch 47/500
 60000/60000 [=====] - 11s 181us/sample - loss: 0.3005 -
 accuracy: 0.8870
 Epoch 48/500
 60000/60000 [=====] - 11s 182us/sample - loss: 0.2984 -
 accuracy: 0.8889
 Epoch 49/500
 60000/60000 [=====] - 11s 181us/sample - loss: 0.2979 -
 accuracy: 0.8892
 Epoch 50/500
 60000/60000 [=====] - 11s 187us/sample - loss: 0.2944 -
 accuracy: 0.8911
 Epoch 51/500
 60000/60000 [=====] - 12s 200us/sample - loss: 0.2948 -
 accuracy: 0.8899
 Epoch 52/500
 60000/60000 [=====] - 12s 193us/sample - loss: 0.2978 -
 accuracy: 0.8906

Epoch 53/500
60000/60000 [=====] - 12s 199us/sample - loss: 0.2960 -
accuracy: 0.8907

Epoch 54/500
60000/60000 [=====] - 12s 192us/sample - loss: 0.2937 -
accuracy: 0.8904

Epoch 55/500
60000/60000 [=====] - 11s 190us/sample - loss: 0.2951 -
accuracy: 0.8906

Epoch 56/500
60000/60000 [=====] - 12s 205us/sample - loss: 0.2934 -
accuracy: 0.8900

Epoch 57/500
60000/60000 [=====] - 11s 190us/sample - loss: 0.2929 -
accuracy: 0.8913

Epoch 58/500
60000/60000 [=====] - 12s 192us/sample - loss: 0.2906 -
accuracy: 0.8924

Epoch 59/500
60000/60000 [=====] - 11s 190us/sample - loss: 0.2905 -
accuracy: 0.8919

Epoch 60/500
60000/60000 [=====] - 12s 193us/sample - loss: 0.2892 -
accuracy: 0.8927

Epoch 61/500
60000/60000 [=====] - 12s 192us/sample - loss: 0.2925 -
accuracy: 0.8912

Epoch 62/500
60000/60000 [=====] - 12s 193us/sample - loss: 0.2868 -
accuracy: 0.8935

Epoch 63/500
60000/60000 [=====] - 12s 192us/sample - loss: 0.2924 -
accuracy: 0.8911

Epoch 64/500
60000/60000 [=====] - 12s 194us/sample - loss: 0.2904 -
accuracy: 0.8925

Epoch 65/500
60000/60000 [=====] - 12s 193us/sample - loss: 0.2887 -
accuracy: 0.8939

Epoch 66/500
60000/60000 [=====] - 12s 195us/sample - loss: 0.2895 -
accuracy: 0.8924

Epoch 67/500
60000/60000 [=====] - 12s 193us/sample - loss: 0.2854 -
accuracy: 0.8929

Epoch 68/500
60000/60000 [=====] - 12s 195us/sample - loss: 0.2819 -
accuracy: 0.8947

```
Epoch 69/500
60000/60000 [=====] - 12s 195us/sample - loss: 0.2869 -
accuracy: 0.8941
Epoch 70/500
60000/60000 [=====] - 12s 198us/sample - loss: 0.2855 -
accuracy: 0.8952
Epoch 71/500
60000/60000 [=====] - 12s 193us/sample - loss: 0.2846 -
accuracy: 0.8937
Epoch 72/500
60000/60000 [=====] - 12s 195us/sample - loss: 0.2845 -
accuracy: 0.8940
Epoch 73/500
60000/60000 [=====] - 12s 194us/sample - loss: 0.2837 -
accuracy: 0.8952
```

[]:

Iterate a few times and find a dropout model that seems to bring the testing and training numbers closer together. When you are done, document your findings in the table below. The ?s are placeholders for accuracy and loss values.

Dropout (Y/N)	Train/Test	Accuracy	Loss
N	Train	.95	.1
N	Test	.87	.5
Y	Train	.89	.28
Y	Test	.87	.35

```
[ ]: train_loss = history.history['loss'][-1]
train_accuracy = history.history['accuracy'][-1]
(test_loss, test_accuracy) = model.evaluate(test_df[FEATURES], test_df[TARGETS])

print('Training loss:', train_loss)
print('Training accuracy:', train_accuracy)

print('Test loss:', test_loss)
print('Test accuracy:', test_accuracy)
```

```
WARNING:tensorflow:Falling back from v2 loop because of error: Failed to find
data adapter that can handle input: <class 'pandas.core.frame.DataFrame'>,
<class 'NoneType'>
10000/10000 [=====] - 1s 68us/sample - loss: 0.3501 -
accuracy: 0.8751
Training loss: 0.2836518274118503
Training accuracy: 0.8952
Test loss: 0.3500618842363358
```


Test accuracy: 0.8751

1.5 Make Predictions

We have now trained the model while trying to reduce overfitting. Let's say we're happy with our numbers and are ready to deploy the model. Now it is time to make predictions.

We could now snap an image of a clothing item, resize it to 28 by 28, and grayscale it. But that is a lot of work and outside the scope of this class. For simplicity, let's use the test images as input to the model and see what predictions we get.

We'll use the `model.predict()` function to do this. Let's make our predictions and peek at the first result.

```
[ ]: predictions = model.predict(test_df[FEATURES])

predictions[0]
```

```
WARNING:tensorflow:Falling back from v2 loop because of error: Failed to find
data adapter that can handle input: <class 'pandas.core.frame.DataFrame'>,
<class 'NoneType'>
```

```
[ ]: array([3.2000608e-10, 5.1018556e-12, 1.2096562e-09, 1.9276606e-11,
          2.5907356e-11, 1.1488260e-03, 9.3423907e-11, 3.7617382e-02,
          1.1600423e-09, 9.6123374e-01], dtype=float32)
```

What are those numbers?

For each image: * the prediction result is in the form of 10 numbers, one for each possible label
* each number represents the level of confidence that a label is the correct label for the particular image
* all 10 numbers should add up to the sum of 1

Let's see if that is true.

```
[ ]:

[ ]: sum(predictions[0]), sum(predictions[1])

[ ]: (0.9999999455570759, 0.9999999934343045)
```

Well, maybe not 1, but the result definitely approaches 1. Floating point math makes summing to exactly 1 a little difficult.

Let's find out which label has the highest predicted number and whether it matches with the actual test label.

To find the highest predicted number we will use Numpy's `argmax` function which returns the index of the maximum value in an array.

```
[ ]: import numpy as np

print('Label with the highest confidence: {predicted_label}'.format(
    predicted_label = np.argmax(predictions[0])))
```

```
print('Actual label: {actual_label}'.format(actual_label = test_labels[0]))
```

Label with the highest confidence: 9

Actual label: 9

With our model the predicted class was class 9, and the actual class was class 9. Success!

1.5.1 Exercise 6: Thresholds

When making our predictions, we blindly accepted the output of `argmax` without really understanding what `argmax` was doing.

`argmax` returns the index of the maximum value in an array. What if there are ties? What happens for a 10-element array that looks like:

```
[0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1]
```

In this case it is a virtual tie between all of the classes. `argmax` will return the first value in the case of a tie. This is problematic for a few reasons. In this case we clearly have little confidence in any class, yet an algorithm that relies on `argmax` would naively predict the first class.

For this exercise, discuss ways we can get around relying solely on `argmax`. Are there better ways of finding a prediction algorithm?

Student Solution

Argmax is good in cases where classes of labels are discrete and the numerical difference between these values are more apparent. This was certainly the case for the model used, where integer values such as 1 and 2 were used to classify certain labels. In cases where this is not apparent, more concrete algorithms that can sort these values and return a maximum value would be better to use

1.6 Exercise 7: MNIST Digits

Another popular MNIST dataset is the [digits dataset](#). This dataset consists of images of labelled, hand-written digits ranging from 0 through 9.

In this exercise you will build a model that predicts the class of MNIST digit images.

The dataset is part of scikit-learn.

```
[ ]: from sklearn import datasets

import pandas as pd

digits_bunch = datasets.load_digits()
digits = pd.DataFrame(digits_bunch.data)
digits['digit'] = digits_bunch.target

#targets
```

```
digits['digit']
```

```
#digits.describe()
```

```
[ ]: 0      0
      1      1
      2      2
      3      3
      4      4
      ..
     1792    9
     1793    0
     1794    8
     1795    9
     1796    8
Name: digit, Length: 1797, dtype: int64
```

You will need to:

- Perform EDA on the data
- Choose a model (or models) to use to predict digits
- Perform any model-specific data manipulation
- Train the model and, if possible, visualize training progression
- Perform a final test of the model on holdout data

Use as many code and text cells as you need to. Explain your work.

1.6.1 Student Solution

EDA

```
[ ]: FEATURES_2 = train_df.columns[:-1]

#finding target classifications
print(digits['digit'].unique())

#loading original dataframe
og_df = digits

og_df
```

```
[0 1 2 3 4 5 6 7 8 9]
```

```
[ ]:      0      1      2      3      4      5      6      7      8      9  ...  55  56  \
0      0.0  0.0   5.0  13.0   9.0   1.0  0.0  0.0  0.0  0.0  ...  0.0  0.0
1      0.0  0.0   0.0  12.0  13.0   5.0  0.0  0.0  0.0  0.0  ...  0.0  0.0
```

2	0.0	0.0	0.0	4.0	15.0	12.0	0.0	0.0	0.0	0.0	...	0.0	0.0
3	0.0	0.0	7.0	15.0	13.0	1.0	0.0	0.0	0.0	8.0	...	0.0	0.0
4	0.0	0.0	0.0	1.0	11.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
...
1792	0.0	0.0	4.0	10.0	13.0	6.0	0.0	0.0	0.0	1.0	...	0.0	0.0
1793	0.0	0.0	6.0	16.0	13.0	11.0	1.0	0.0	0.0	0.0	...	0.0	0.0
1794	0.0	0.0	1.0	11.0	15.0	1.0	0.0	0.0	0.0	0.0	...	0.0	0.0
1795	0.0	0.0	2.0	10.0	7.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0
1796	0.0	0.0	10.0	14.0	8.0	1.0	0.0	0.0	0.0	2.0	...	0.0	0.0

	57	58	59	60	61	62	63	digit
0	0.0	6.0	13.0	10.0	0.0	0.0	0.0	0
1	0.0	0.0	11.0	16.0	10.0	0.0	0.0	1
2	0.0	0.0	3.0	11.0	16.0	9.0	0.0	2
3	0.0	7.0	13.0	13.0	9.0	0.0	0.0	3
4	0.0	0.0	2.0	16.0	4.0	0.0	0.0	4
...
1792	0.0	2.0	14.0	15.0	9.0	0.0	0.0	9
1793	0.0	6.0	16.0	14.0	6.0	0.0	0.0	0
1794	0.0	2.0	9.0	13.0	6.0	0.0	0.0	8
1795	0.0	5.0	12.0	16.0	12.0	0.0	0.0	9
1796	1.0	8.0	12.0	14.0	12.0	1.0	0.0	8

[1797 rows x 65 columns]

```
[ ]: #One hot encoding

target_columns =
    ↳ ['target_0', 'target_1', 'target_2', 'target_3', 'target_4', 'target_5', 'target_6', 'target_7', 't
TARGETS = target_columns
ohe = OneHotEncoder(sparse=False)

encodedTargets = ohe.fit_transform(og_df[['digit']])

encodedTargets_df = pd.DataFrame(encodedTargets)
encodedTargets_df.columns =target_columns

og_df = og_df.drop('digit',axis =1 )

og_df = og_df.join(encodedTargets_df)

[ ]: og_df
```

```
[ ]:      0      1      2      3      4      5      6      7      8      9  ...  target_0  \
0      0.0  0.0   5.0  13.0   9.0   1.0  0.0  0.0  0.0  0.0  ...      1.0
1      0.0  0.0   0.0  12.0  13.0   5.0  0.0  0.0  0.0  0.0  ...      0.0
2      0.0  0.0   0.0   4.0  15.0  12.0  0.0  0.0  0.0  0.0  ...      0.0
3      0.0  0.0   7.0  15.0  13.0   1.0  0.0  0.0  0.0  8.0  ...      0.0
4      0.0  0.0   0.0   1.0  11.0   0.0  0.0  0.0  0.0  0.0  ...      0.0

... ..
1792  0.0  0.0   4.0  10.0  13.0   6.0  0.0  0.0  0.0  1.0  ...      0.0
1793  0.0  0.0   6.0  16.0  13.0  11.0  1.0  0.0  0.0  0.0  ...      1.0
1794  0.0  0.0   1.0  11.0  15.0   1.0  0.0  0.0  0.0  0.0  ...      0.0
1795  0.0  0.0   2.0  10.0   7.0   0.0  0.0  0.0  0.0  0.0  ...      0.0
1796  0.0  0.0  10.0  14.0   8.0   1.0  0.0  0.0  0.0  2.0  ...      0.0

      target_1  target_2  target_3  target_4  target_5  target_6  target_7  \
0          0.0        0.0        0.0        0.0        0.0        0.0        0.0
1          1.0        0.0        0.0        0.0        0.0        0.0        0.0
2          0.0        1.0        0.0        0.0        0.0        0.0        0.0
3          0.0        0.0        1.0        0.0        0.0        0.0        0.0
4          0.0        0.0        0.0        1.0        0.0        0.0        0.0

... ..
1792        0.0        0.0        0.0        0.0        0.0        0.0        0.0
1793        0.0        0.0        0.0        0.0        0.0        0.0        0.0
1794        0.0        0.0        0.0        0.0        0.0        0.0        0.0
1795        0.0        0.0        0.0        0.0        0.0        0.0        0.0
1796        0.0        0.0        0.0        0.0        0.0        0.0        0.0

      target_8  target_9
0          0.0        0.0
1          0.0        0.0
2          0.0        0.0
3          0.0        0.0
4          0.0        0.0

... ..
1792        0.0        1.0
1793        0.0        0.0
1794        1.0        0.0
1795        0.0        1.0
1796        1.0        0.0

[1797 rows x 74 columns]
```

Splitting the data

```
[ ]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(
    og_df[FEATURES_2],
    og_df[TARGETS],
    test_size=0.2,
```

```

random_state=180, shuffle=True)

train_df = X_train.join(y_train)
test_df = X_test.join(y_test)

```

Building the Model

Training set

```

[ ]: model = tf.keras.Sequential([
    tf.keras.layers.Dense(128, input_shape=(len(FEATURES_2),)),
    tf.keras.layers.Dropout(rate=0.25),
    tf.keras.layers.Dense(64, activation=tf.nn.relu),
    tf.keras.layers.Dropout(rate=0.15),
    tf.keras.layers.Dense(32, activation=tf.nn.relu),
    tf.keras.layers.Dropout(rate=0.01),
    tf.keras.layers.Dense(len(TARGETS), activation=tf.nn.softmax)
])

model.compile(
    loss='categorical_crossentropy',
    optimizer='Adam',
    metrics=['accuracy'],
)

callback = tf.keras.callbacks.EarlyStopping(monitor='loss', patience=5)

history = model.fit(
    train_df[FEATURES_2],
    train_df[TARGETS],
    epochs=500,
    callbacks=[callback]
)

```

WARNING:tensorflow:Falling back from v2 loop because of error: Failed to find data adapter that can handle input: <class 'pandas.core.frame.DataFrame'>, <class 'NoneType'>

Train on 1437 samples

Epoch 1/500

1437/1437 [=====] - 1s 507us/sample - loss: 3.0616 - accuracy: 0.2262

Epoch 2/500

1437/1437 [=====] - 0s 152us/sample - loss: 1.3810 - accuracy: 0.5463

Epoch 3/500

1437/1437 [=====] - 0s 151us/sample - loss: 0.9188 -
 accuracy: 0.7119
 Epoch 4/500
 1437/1437 [=====] - 0s 154us/sample - loss: 0.6370 -
 accuracy: 0.7989
 Epoch 5/500
 1437/1437 [=====] - 0s 156us/sample - loss: 0.4558 -
 accuracy: 0.8497
 Epoch 6/500
 1437/1437 [=====] - 0s 173us/sample - loss: 0.3616 -
 accuracy: 0.8817
 Epoch 7/500
 1437/1437 [=====] - 0s 165us/sample - loss: 0.2692 -
 accuracy: 0.9172
 Epoch 8/500
 1437/1437 [=====] - 0s 154us/sample - loss: 0.3099 -
 accuracy: 0.9005
 Epoch 9/500
 1437/1437 [=====] - 0s 153us/sample - loss: 0.2069 -
 accuracy: 0.9318
 Epoch 10/500
 1437/1437 [=====] - 0s 153us/sample - loss: 0.1674 -
 accuracy: 0.9429
 Epoch 11/500
 1437/1437 [=====] - 0s 155us/sample - loss: 0.2061 -
 accuracy: 0.9297
 Epoch 12/500
 1437/1437 [=====] - 0s 154us/sample - loss: 0.1802 -
 accuracy: 0.9402
 Epoch 13/500
 1437/1437 [=====] - 0s 152us/sample - loss: 0.1666 -
 accuracy: 0.9492
 Epoch 14/500
 1437/1437 [=====] - 0s 153us/sample - loss: 0.1578 -
 accuracy: 0.9499
 Epoch 15/500
 1437/1437 [=====] - 0s 153us/sample - loss: 0.1355 -
 accuracy: 0.9527
 Epoch 16/500
 1437/1437 [=====] - 0s 152us/sample - loss: 0.1381 -
 accuracy: 0.9589
 Epoch 17/500
 1437/1437 [=====] - 0s 152us/sample - loss: 0.1181 -
 accuracy: 0.9610
 Epoch 18/500
 1437/1437 [=====] - 0s 153us/sample - loss: 0.1287 -
 accuracy: 0.9624
 Epoch 19/500

1437/1437 [=====] - 0s 152us/sample - loss: 0.1133 -
 accuracy: 0.9589
 Epoch 20/500
 1437/1437 [=====] - 0s 150us/sample - loss: 0.1016 -
 accuracy: 0.9624
 Epoch 21/500
 1437/1437 [=====] - 0s 151us/sample - loss: 0.1066 -
 accuracy: 0.9666
 Epoch 22/500
 1437/1437 [=====] - 0s 162us/sample - loss: 0.1106 -
 accuracy: 0.9603
 Epoch 23/500
 1437/1437 [=====] - 0s 152us/sample - loss: 0.0919 -
 accuracy: 0.9617
 Epoch 24/500
 1437/1437 [=====] - 0s 154us/sample - loss: 0.0720 -
 accuracy: 0.9791
 Epoch 25/500
 1437/1437 [=====] - 0s 154us/sample - loss: 0.0749 -
 accuracy: 0.9749
 Epoch 26/500
 1437/1437 [=====] - 0s 152us/sample - loss: 0.0693 -
 accuracy: 0.9777
 Epoch 27/500
 1437/1437 [=====] - 0s 153us/sample - loss: 0.0783 -
 accuracy: 0.9777
 Epoch 28/500
 1437/1437 [=====] - 0s 152us/sample - loss: 0.0818 -
 accuracy: 0.9756
 Epoch 29/500
 1437/1437 [=====] - 0s 157us/sample - loss: 0.0627 -
 accuracy: 0.9798
 Epoch 30/500
 1437/1437 [=====] - 0s 158us/sample - loss: 0.0723 -
 accuracy: 0.9722
 Epoch 31/500
 1437/1437 [=====] - 0s 153us/sample - loss: 0.0748 -
 accuracy: 0.9749
 Epoch 32/500
 1437/1437 [=====] - 0s 151us/sample - loss: 0.0811 -
 accuracy: 0.9701
 Epoch 33/500
 1437/1437 [=====] - 0s 154us/sample - loss: 0.0605 -
 accuracy: 0.9777
 Epoch 34/500
 1437/1437 [=====] - 0s 152us/sample - loss: 0.0430 -
 accuracy: 0.9868
 Epoch 35/500


```

1437/1437 [=====] - 0s 155us/sample - loss: 0.0841 -
accuracy: 0.9756
Epoch 36/500
1437/1437 [=====] - 0s 153us/sample - loss: 0.0596 -
accuracy: 0.9777
Epoch 37/500
1437/1437 [=====] - 0s 153us/sample - loss: 0.0488 -
accuracy: 0.9854
Epoch 38/500
1437/1437 [=====] - 0s 153us/sample - loss: 0.0589 -
accuracy: 0.9833
Epoch 39/500
1437/1437 [=====] - 0s 152us/sample - loss: 0.0616 -
accuracy: 0.9819

```

Graphing Model Process

```

[ ]: plt.figure(figsize=(16,5))

plt.subplot(1,2,2)
plt.plot(history.history['loss'])
plt.title('Training Loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train_loss'], loc='best')


plt.subplot(1,2,1)
plt.plot(history.history['accuracy'])
plt.title('Training Accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train_accuracy'], loc='best')

```

```

[ ]: <matplotlib.legend.Legend at 0x7f954b443e10>

```



Model Evaluation

```
[ ]: train_loss = history.history['loss'][-1]
train_accuracy = history.history['accuracy'][-1]
(test_loss, test_accuracy) = model.evaluate(test_df[FEATURES_2],
↳test_df[TARGETS])

print('Training loss:', train_loss)
print('Training accuracy:', train_accuracy)

print('Test loss:', test_loss)
print('Test accuracy:', test_accuracy)
```

```
WARNING:tensorflow:Falling back from v2 loop because of error: Failed to find
data adapter that can handle input: <class 'pandas.core.frame.DataFrame'>,
<class 'NoneType'>
360/360 [=====] - 0s 553us/sample - loss: 0.0862 -
accuracy: 0.9861
Training loss: 0.061573891125997905
Training accuracy: 0.9819068
Test loss: 0.08623956021102559
Test accuracy: 0.9861111
```

```
[ ]: predictions = model.predict(test_df[FEATURES_2])

predictions[0]
```

```
WARNING:tensorflow:Falling back from v2 loop because of error: Failed to find
data adapter that can handle input: <class 'pandas.core.frame.DataFrame'>,
<class 'NoneType'>
```

```
[ ]: array([1.6430340e-08, 3.3618651e-08, 3.1002009e-14, 1.9662133e-11,
1.2574532e-08, 5.4443963e-11, 1.0000000e+00, 3.1500911e-14,
1.1559946e-08, 1.5749539e-15], dtype=float32)
```